



Keras is an open source deep-learning API for *Python*

Original author - François Chollet

Today - various developers

including *Google's TensorFlow* team (since 2017)

Initial release - 27 March 2015

Stable release - 2.0.8 / 24 August 2017

Optional backends:

TensorFlow, CNTK, Theano



Motivation

Before:

- Undedicated libraries – self implementation
- Long unreadable code, multi-file (*Caffe* for example)
- Limited wrappers like *Lasagne*

With Keras:

- Simple syntax & full control, one *.py run file
- Includes all SOT NN layers & functions
- building complex graphs with ease
- growing community of contributors
(academia & industry)



Getting started

Documentation:
keras.io

Blog:
blog.keras.io

Installation: (or directly from git)
\$ sudo pip install keras
(after pre-requisites & backend)

CUDA is supported by backends, and will be used if installed,
unless flagged out.



Getting started

```
from keras.models import Sequential
from keras.layers import Dense, Activation

# sequential API toy model:
model = Sequential()
model.add(Dense(units=64, input_dim=100))
model.add(Activation('relu'))
model.add(Dense(units=10))
model.add(Activation('softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.SGD(lr=0.01, momentum=0.9, nesterov=True))
model.fit(x_train, y_train, epochs=5, batch_size=32)

# evaluating a model:
loss_and_metrics = model.evaluate(x_test, y_test, batch_size=128)

# model production use:
classes = model.predict(x_test, batch_size=128)
```



Two APIs in Keras

Sequential API - VGG16 example:

```
model = Sequential()
model.add(Convolution2D(64, 3, 3, activation='relu', padding='same', input_shape=(3, 224, 224)))
model.add(Convolution2D(64, 3, 3, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2), strides=(2, 2)))

model.add(Convolution2D(128, 3, 3, activation='relu', padding='same'))
model.add(Convolution2D(128, 3, 3, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2), strides=(2, 2)))

model.add(Convolution2D(256, 3, 3, activation='relu', padding='same'))
model.add(Convolution2D(256, 3, 3, activation='relu', padding='same'))
model.add(Convolution2D(256, 3, 3, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2), strides=(2, 2)))

model.add(Convolution2D(512, 3, 3, activation='relu', padding='same'))
model.add(Convolution2D(512, 3, 3, activation='relu', padding='same'))
model.add(Convolution2D(512, 3, 3, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2), strides=(2, 2)))

model.add(Convolution2D(512, 3, 3, activation='relu', padding='same'))
model.add(Convolution2D(512, 3, 3, activation='relu', padding='same'))
model.add(Convolution2D(512, 3, 3, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2), strides=(2, 2)))

model.add(Flatten())
model.add(Dense(4096, activation='relu'))
model.add(Dense(4096, activation='relu'))
model.add(Dense(1000, activation='softmax'))
```



Two APIs in Keras

Functional API - VGG16 example:

```
img_input = Input(shape=input_shape)
x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv1')(img_input)
x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv2')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block1_pool')(x)

x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv1')(x)
x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv2')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block2_pool')(x)

x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv1')(x)
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv2')(x)
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block3_pool')(x)

x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block4_pool')(x)

x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block5_pool')(x)

x = Flatten(name='flatten')(x)
x = Dense(4096, activation='relu', name='fc1')(x)
x = Dense(4096, activation='relu', name='fc2')(x)
x = Dense(classes, activation='softmax', name='predictions')(x)
```



Configuration files

Configuration file: `~/.keras/keras.json`

```
{  
  "image_data_format": "channels_last",  
  "epsilon": 1e-07,  
  "floatx": "float32",  
  "backend": "tensorflow"  
}
```



Configuration files

When using *Theano* backend: `~/.theanorc`

```
[blas]  
ldflags =
```

```
[global]  
floatX = float32  
device = gpu
```

```
[nvcc]  
fastmath = True
```

```
[gcc]  
cxxflags = -ID:\MinGW\include
```

```
[cuda]  
root=/usr/local/cuda-8.0/
```

Documentation at: deeplearning.net/software/theano/library/config.html#




Basic layer types

- Core Layers – shown in example
- Convolutional Layers – shown in example
- Recurrent Layers – shown in example
- Pooling Layers – shown in example
- Locally-connected Layers
- Embedding Layers (including pretrained)
- Merge Layers – shown in example
- Advanced Activations Layers
- Normalization Layers
- Noise layers
- Layer wrappers
- Writing your own Keras layers

Activation functions

- linear (regression)
- softmax (classification)
- elu
- selu
- softplus
- softsign
- relu
- hard_sigmoid
- sigmoid (logistic function / binary)
- tanh



and other advanced activations as well...



Loss functions

- `mean_squared_error`
- `mean_absolute_error`
- `mean_absolute_percentage_error`
- `mean_squared_logarithmic_error`
- `squared_hinge`
- `hinge`
- `categorical_hinge`
- `logcosh`
- `categorical_crossentropy`
- `sparse_categorical_crossentropy`
- `binary_crossentropy`
- `kullback_leibler_divergence`
- `poisson`
- `cosine_proximity`



Metrics

- `binary_accuracy`
- `categorical_accuracy`
- `sparse_categorical_accuracy`
- `top_k_categorical_accuracy`
- `sparse_top_k_categorical_accuracy`
- custom metrics (user defined)



Optimizers

- SGD
- RMSprop
- Adagrad
- Arguments
- References
- Adadelta optimizer.
- Arguments
- References
- Adam
- Adamax
- Nadam
- TFOptimizer - TensorFlow's



Basic layer types

- Core Layers – shown in example
- Convolutional Layers – shown in example
- Recurrent Layers – shown in example
- Pooling Layers – shown in example
- Locally-connected Layers
- Embedding Layers (including pre-trained)
- Merge Layers – shown in example
- Advanced Activations Layers
- Normalization Layers
- Noise layers
- Layer wrappers
- Writing your own Keras layers



Customizing

Lambda layer – untrainable

(e.g. custom activation function)

Writing your own layer – trainable

Writing your own metrics

Using your own generator as data flow

for pre-processing on-the-fly (during training)

Scikit-learn API for embedding a mixed model



Showcase – stocks

Daily values vs. DJi

